



Hamiltonian Monte Carlo with Energy Conserving Subsampling (HMC-EC)

Markov Chain Monte Carlo - Theory and practical applications

Ancellin Eloi - Eisenbarth Florian - Foucault Armand - Joël Garde

- Introduction
- Bayesian context
- Subsampling
- HMC
- Pseudo-marginal
- HMC-with-Gibbs
- Estimators
- Experimental results

Introduction

- Bayesian Context:
 - Dataset $D \in \mathbb{R}^{n \times d}$
 - Points coming from law that we do not know how to sample from
- Objective
 - **HMC-within gibbs** to sample from a posterior distribution
 - 2 major issues:
 - 1) Explore the posterior distribution efficiently when **d** is large  HMC methods
 - 2) Keep a reasonable computation cost when **n** is large  Subsampling

1) and 2) at the same time is harder!

Subsampling

- Advantages

- Driving cost in HMC : evaluation of both the **likelihood** and its **gradient** $\Rightarrow O(n)$
- Subsample of size **m** \Rightarrow accelerate computation of HMC
- $O(m) \ll O(n)$

- Issues:

- Naive subsampling \Rightarrow Disconnects the Hamiltonien from its own dynamics \Rightarrow Deterioration of energy conserving property \Rightarrow Loses ability to sample efficiently in high dimension

How can we solve problems 1) and 2) at the same time ?

Bayesian inference

- Dataset: $D \in R^{n \times d}$, labels $Y \in [-1, 1]^n$
- Model $\Phi_\theta(x) = \hat{y}$
- Posterior distribution $\pi(\theta | x)$
- Likelihood $L(\theta) = P(x | \theta)$
- Prior $\pi(\theta)$
- Bayes Theorem
$$\pi(\theta | x) = \frac{\pi(\theta)P(x | \theta)}{\int \pi(u)P(x | u)du}$$
$$\pi(\theta | x) \propto \pi(\theta)P(x|\theta)$$

Hamiltonian Monte Carlo (HMC) (1/2)

- Position: $\theta \in \mathbb{R}^d$
- Target distribution: $\pi(\theta) \propto e^{-U(\theta)}$
- Momentum: $p \in \mathbb{R}^d$
- Potential energy: $U(\theta) = \ell(\theta)$ where ℓ is the log-likelihood
- Kinetic energy: $K(p) = \frac{p^T M^{-1} p}{2}$
- Extended target density: $\Pi(\theta, p) \propto \exp\{-U(\theta) - K(p)\} \quad p, q \in \mathbb{R}^d$
- The Hamiltonian: $H(\theta, p) = U(\theta) + K(p)$

Hamiltonian Monte Carlo (HMC) (2/2)

- Hamiltonian dynamics:

$$\begin{aligned}\frac{\partial H}{\partial \theta^{t,i}}(\theta^t, p^t) &= \frac{\partial U(\theta^t)}{\partial \theta^{t,i}} = -\frac{dp^{t,i}}{dt} \\ \frac{\partial H}{\partial p^{t,i}}(\theta^t, p^t) &= \frac{\partial K(p^t)}{\partial p^{t,i}} = p^{t,i} = \frac{d\theta^{t,i}}{dt}\end{aligned}$$

- Leapfrog integrator:

$$\begin{aligned}p^{k+\frac{1}{2}} &= p^k - \left(\frac{h}{2}\right)\nabla U(\theta^k) \\ \theta^{k+1} &= \theta^k + hM^{-1}p^{k+\frac{1}{2}} \\ p^{k+1} &= p^{k+\frac{1}{2}} - \left(\frac{h}{2}\right)\nabla U(\theta^{k+1})\end{aligned}$$

Pseudo-Marginal MCMC (1/2)

- Objective: sample from an analytically unknown target distribution $\pi(\theta)$
- Idea: use an unbiased and non-negative estimator :

$$\bar{\pi}(\theta, u) = \bar{\pi}(\theta|u) p_U(u)$$

Where u is an auxiliary variable with distribution p_U

$$\Rightarrow \bar{\pi} \text{ verifies } E_{p_U}[\bar{\pi}(\theta, u)] = \pi(\theta)$$

Pseudo-Marginal MCMC (2/2)

- In our case, we consider a Bayesian setting, therefore:

$$\bar{\pi}(\theta, u) \propto \hat{L}_m(\theta) p_{\Theta}(\theta) p_U(u)$$

Where u is a subset of our observations and $\hat{L}_m(\theta)$ is an estimator of the likelihood of θ based on the subset u . $p_{\Theta}(\theta)$ is the prior on θ .

➡ Use the above extended density as the target of a MH algorithm in order to obtain samples from π

Hamiltonian Monte Carlo Within Gibbs (1/6)

A “Gibbs framework”, for each iteration, makes alternatively:

- A subsampling step: draw \mathbf{u} among your set of observations \mathbf{y}
- An HMC step: make a proposal for your new parameter value, according to the Hamiltonian dynamics

Hamiltonian Monte Carlo Within Gibbs (2/6)

$$\bar{\pi}(\theta, p, u) \propto \exp(-\hat{\mathcal{H}}(\theta, p)) P_U(u)$$

$$\begin{aligned} u_{k+1} &\sim u \mid \theta_k, p_k \\ \theta_{k+1}, p_{k+1} &\sim \theta, p \mid u_{k+1} \end{aligned}$$

1/ Subsample auxiliary variables

- Size m of the subsample chosen such that the MCMC walk is computationally efficient:
 - m too small \rightarrow highly unstable estimator (high variance) \rightarrow low acceptance rate \rightarrow waste of computational resources
 - m too large \rightarrow too many computations \rightarrow too much computational time

$$\Rightarrow m = n^{0.5}$$

Hamiltonian Monte Carlo Within Gibbs (4/6)

1/ Subsample auxiliary variables

- Induce correlation within the draws to tolerate higher variance in the estimator
 - By inducing correlation between the u 's, we can tolerate a substantially smaller m

➡ Divide u into G blocks and update only one block at a time randomly ➡ $n = n^{0.5}$

- Propose u' according to p_U and accept it with probability

$$\min \left(1, \frac{\hat{L}_m(\theta_{t-1}, u')}{\hat{L}_m(\theta_{t-1}, u_{t-1})} \right)$$

Hamiltonian Monte Carlo Within Gibbs (5/6)

2/ HMC step

- Recall the estimated Hamiltonian:

$$\hat{H}(\theta, p) := \hat{\mathcal{U}}(\theta) + \mathcal{K}(p)$$

- We simulate the Hamiltonian dynamics by leapfrog discretization for L steps per HMC step

➡ We start at $(\theta_0, p_0) = (\theta_{t-1}, p_0)$ and arrive at proposal $(\theta^L, -p^L)$

Hamiltonian Monte Carlo Within Gibbs (6/6)

2/ HMC step

- Accept the proposal with probability

$$\min \left(1, \exp \left(-\hat{H}(\theta_L, -p_L) + \hat{H}(\theta_0, p_0) \right) \right)$$

Estimator of the likelihood

	Perturbed	Signed
Unbiased	×	✓
Positive	✓	×
Cheap to compute	✓	✓

Difference estimator

$$\hat{h}(\theta) = \sum_{i=0}^n q_i(\theta) + \frac{m}{n} \sum_{u_i \in u} \ell_{u_i}(\theta) - q_{ui}(\theta)$$



2nd order
approximation of the full
likelihood



Correction based on
batch u

Other Tricks

- Dual averaging:
 - Step size automatic.
 - Next step: full blown no-U-turn
- Approximated optimal Euclidean-Gaussian kinetic energies


$$p \sim \mathcal{N}(0, M_{\theta^*})$$

$$M_{\theta^*} \approx -H_L(\theta^*) + s^2 I_d$$


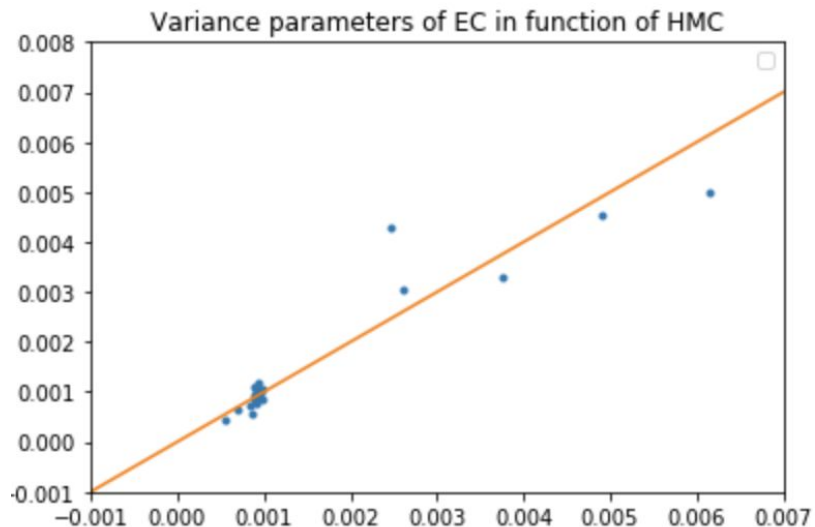
Experimental Results (1/3)

Obtained parameters distribution of HMC(Black) and EC(Blue) for 20 000 iterations.

samples repartitions and logistic regression coeffs

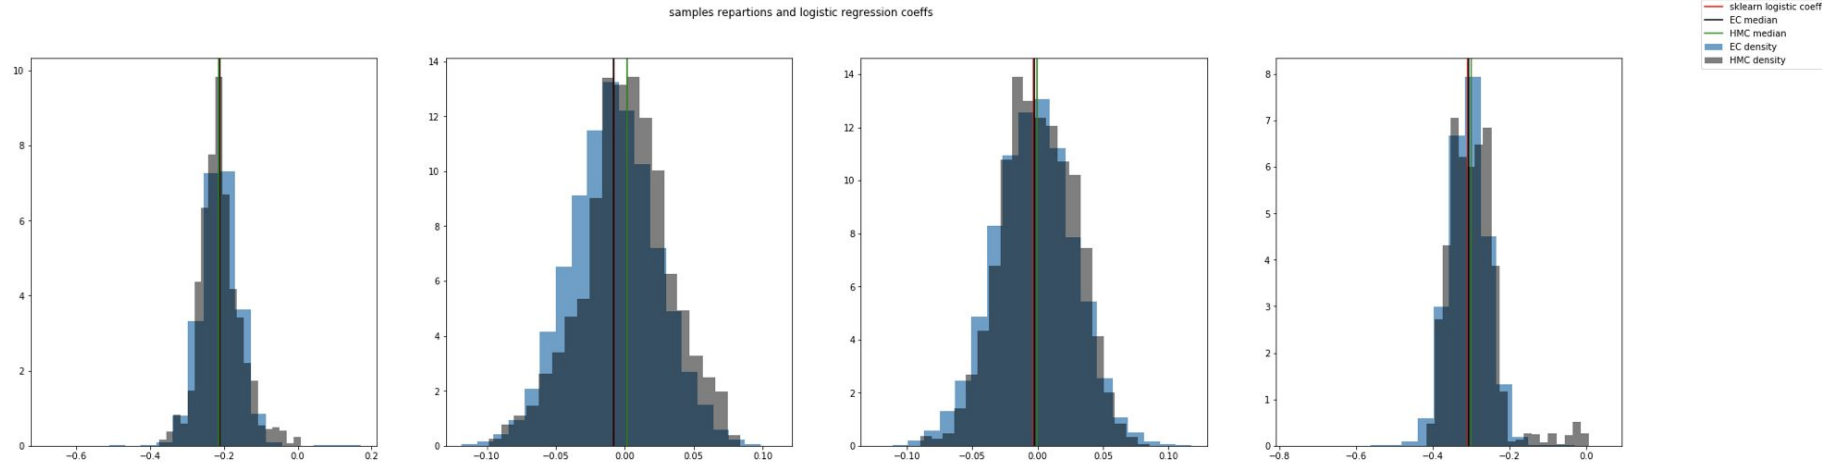


Experimental Results (2/3)



Experimental Results (3/3)

Running the algorithms with 4 000 iterations for the EC and 300 for the HMC (to have a comparable runtime)



Conclusion

Experimental results prove that:

- Inefficiency factor for HMC = 2.08 for EC = 2.185
- Need 5% more iterations for EC than HMC
- EC is 30 times faster than HMC in our measurements (the gain is \sqrt{n})

Left to explore:

- No-U-turn to have the good number of leapfrog steps automatically
- Parameter tuning
- Include practical solutions into a theoretical framework (ex: Stepsize lock)
- Parallelize MCMC to fasten the calculations

Code :

- HMC Energy conservative Kernel
- Leapfrog integrator
- Block poisson
- Log likelihood

```

1 def hmc_ec_kernel(u, q, p_est, e, L, Hs, i, G, x, y, lmbda, shrink, log_delta, t0, nu, gamma, kappa):
2
3     #pseudo-step
4     u_prime = block_update(u.copy(), G, x)
5     log_uniform = np.log(np.random.uniform(0,1))
6     log_alpha_pseudo = hat_diff(u_prime, q, *p_est, x, y) - hat_diff(u, q, *p_est, x, y)
7     do_accept = log_uniform < log_alpha_pseudo
8     if do_accept:
9         u = u_prime
10
11     #hamiltonian step
12     p = multivariate_normal(L) #to do: optimize the inversion away.
13     qL, pL = leapfrog(q.copy(), p.copy(), u, p_est, e, np.minimum(int(lmbda/e)+1, 10), L, shrink, x, y)
14     log_uniform = np.log(np.random.uniform(0,1))
15     log_alpha_hmc = - H(qL, pL, u, p_est, shrink, L, x, y) + H(q, p, u, p_est, shrink, L, x, y)
16     do_accept = log_uniform < log_alpha_hmc
17     if do_accept:
18         q = qL
19
20     #dual averaging (not done for the moment, some nans errors) with some difference (log / exp)
21     #exponential averaging of the errors
22     #as to be re-initialized each time we update H.
23     Hs_next = np.exp(log_delta) - np.minimum(np.exp(log_alpha_hmc), 1.0)
24     Hs = (1 - (1/(i + t0))) * Hs + (1/(i + t0)) * Hs_next
25     log_e = nu - (i**(0.5) / gamma) * Hs
26     e_next = i**(-kappa) * log_e + (1 - i**(-kappa)) * np.log(e)
27     e = np.exp(e_next)
28     return u, q, e, Hs, i+1.0, log_alpha_hmc
29

```



```

1 @njit(fastmath=True)
2 def leapfrog(q, p, u, p_est, e, D, L, shrink, x, y):
3
4     """
5     Leapfrog approximator, step size e, length D*e.
6     q: cinetic variable.
7     p: momemtum variable.
8     L: cholesky lower part of M
9     grad_U: gradient of the cinetic energy.
10    """
11    #M_i = np.linalg.inv(L.T @ L) weird, it should be L.T @ L to give me M.
12    for i in range(D):
13        p -= 0.5 * e * grad_U(q, p, u, p_est, shrink, x, y)
14        q += e * np.linalg.solve(L.T, np.linalg.solve(L, p)) #should be doing a fast triangular solve
        instead. #or just get the inverse since the matrix is small
15        p -= 0.5 * e * grad_U(q, p, u, p_est, shrink, x, y)
16    return q, -p
17

```

```

1 def log_block_poisson(u, lmbda, a, l, q, qs, prim, grad, hess):
2     """ logarithm of block-poisson estimator.
3     expects u to be a list of list of vectors of indices.(lambda * xi_lambda)
4     """
5     qk_sum = control_v(q, qs, prim, grad, hess)
6     sum_log_diff = 0
7     for ul in u:
8         for ui in ul:
9             diff = log_likelihood(y[ui], x[ui], q) - control_v(q, *precompute(y[ui], x[ui], qs))
10            sum_log_diff += np.log(np.abs((diff - a) / lmbda))
11
12    return qk_sum + (a + lmbda) + sum_log_diff
13
14
15 def grad_log_block_poisson(u, lmbda, a, l, q, qs, prim, grad, hess):
16     """ gradient of the log_block_poisson estimator.
17     """
18     n = x.shape[0]
19     grad_qk_sum = grad + hess @ (q - qs)
20     grad_sum_log_diff = 0.0
21     for ul in u:
22         for ui in ul:
23             qs, v, grad_batch, hess_batch = precompute(y[ui], x[ui], qs)
24             diff = log_likelihood(y[ui], x[ui], q) - control_v(q, qs, v, grad_batch, hess_batch)
25             grad_control_batch = grad_batch + hess_batch @ (q - qs)
26             grad_diff_batch = grad_log_likelihood(y[ui], x[ui], q) - grad_control_batch
27             grad_sum_log_diff += grad_diff_batch / (diff - a) * lmbda
28    return grad_qk_sum + grad_sum_log_diff
29
30

```

```

30
31 @njit(fastmath=True)
32 def block_update(u, G, x):
33     """block-update of u considering G blocks.
34     """
35     m, n = u.shape[0], x.shape[0]
36     group_size = m // G
37     replace = np.arange(m) < group_size
38     new_u = np.random.choice(n, u.shape, replace=False)
39     replace = np.random.permutation(replace)
40     return np.where(replace, new_u, u)
41
42 def poisson_update(u, m, lmbda, kappa):
43     """update of u for a poisson law.
44     m is the size of the base batch.
45     only kappa list of batches are updated to correlate u's.
46     correlation ~ 1 - kappa / lmbda
47     returns a list of list of batches. (lambda * xi_lambda batches)
48     """
49     n = x.shape[0]
50     indices_to_update = np.random.choice(lmbda, kappa, replace=False)
51
52     xi = np.random.poisson(1, kappa)
53     for idx in indices_to_update:
54         u[idx] = []
55         for j in range(xi):
56             batch = np.random.choice(n, m, replace=False)
57             u[idx].append(batch)
58     return u

```

```

1 @njit(fastmath=True)
2 def precompute(y, x, qs):
3     """ return val, primal, gradient and hessian at point qs
4     """
5     px = (y * 2 - 1).reshape(-1,1) * x
6     u = 1 + np.exp(- px @ qs)
7     v = u ** (-1)
8     return qs, -np.log(u).sum(), (px * (1 - v).reshape(-1,1)).sum(axis=0), - (px.T * (v * (1 - v))) @ px
9
10 @njit(fastmath=True)
11 def hat_diff(u, q, qs, prim, grad, hess, x, y):
12     """log-likelihood difference estimator.
13     """
14     n = x.shape[0]
15     m = u.shape[0]
16     diff = log_likelihood(y[u], x[u], q) - control_v(q, *precompute(y[u], x[u], qs))
17     return control_v(q, qs, prim, grad, hess) + (n / m) * diff
18
19 @njit(fastmath=True)
20 def grad_hat_diff(u, q, qs, prim, grad, hess, x, y):
21     """ Estimator of the gradient.
22     In our case, it is also the gradient of the estimator.
23     Assumes that hessian is a symmetric matrix (which should be the case.)
24     """
25     n = x.shape[0]
26     m = u.shape[0]
27     grad_control_full = grad + hess @ (q - qs)
28     _, _, grad_batch, hess_batch = precompute(y[u], x[u], qs)
29     grad_control_batch = grad_batch + hess_batch @ (q - qs)
30     grad_diff_batch = grad_log_likelihood(y[u], x[u], q) - grad_control_batch
31     return grad_control_full + (n / m) * grad_diff_batch
32
33 @njit(fastmath=True)
34 def control_v(q, qs, prim, grad, hess):
35     """
36     2nd order approximation of the loglikelihood at qs, evaluated at q
37     given the precomputed primal, gradient and hessian.
38     """
39     d = q - qs
40     return prim + grad @ d + 0.5 * d.T @ (hess @ d)
41

```

```

1 def log_likelihood(y, x, q):
2     """ Likelihood for  $y = \{0, 1\}^n$ 
3     """
4     p = y * 2 - 1 #translate y to -1, 1
5     return - np.log(1 + np.exp(- p * (x @ q))).sum()
6
7 @njit(fastmath=True)
8 def grad_log_likelihood(y, x, q):
9     p = y * 2 - 1 #translate y to -1, 1
10    return (x * ( p * (1 + np.exp(p * (x @ q)))** (-1) ).reshape((-1,1))).sum(axis = 0)
11
12 @njit(fastmath=True)
13 def hessian_log_likelihood(y, x, q):
14     p = y * 2 - 1
15     return (p.reshape((-1,1))* x).T @ ( x * (p * ((1 + np.exp( - p * (x @ q)))** (-1))).reshape((-1,1)))

```